


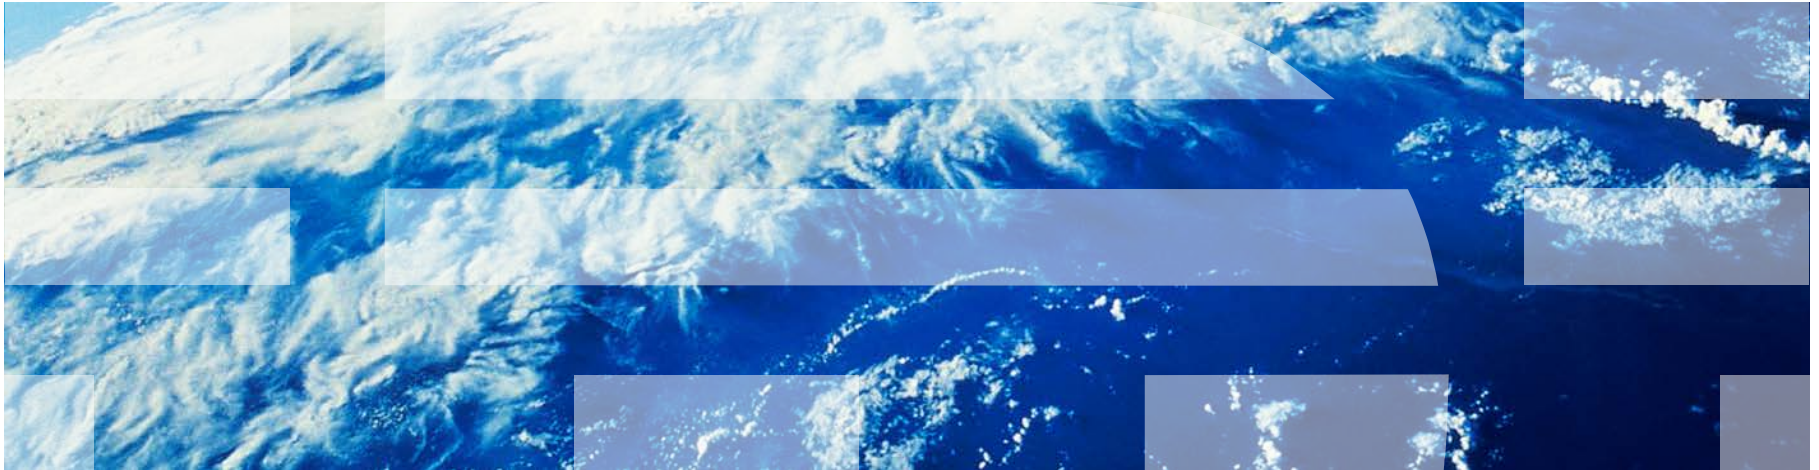
Computer Systems for Data Science

Topic 3

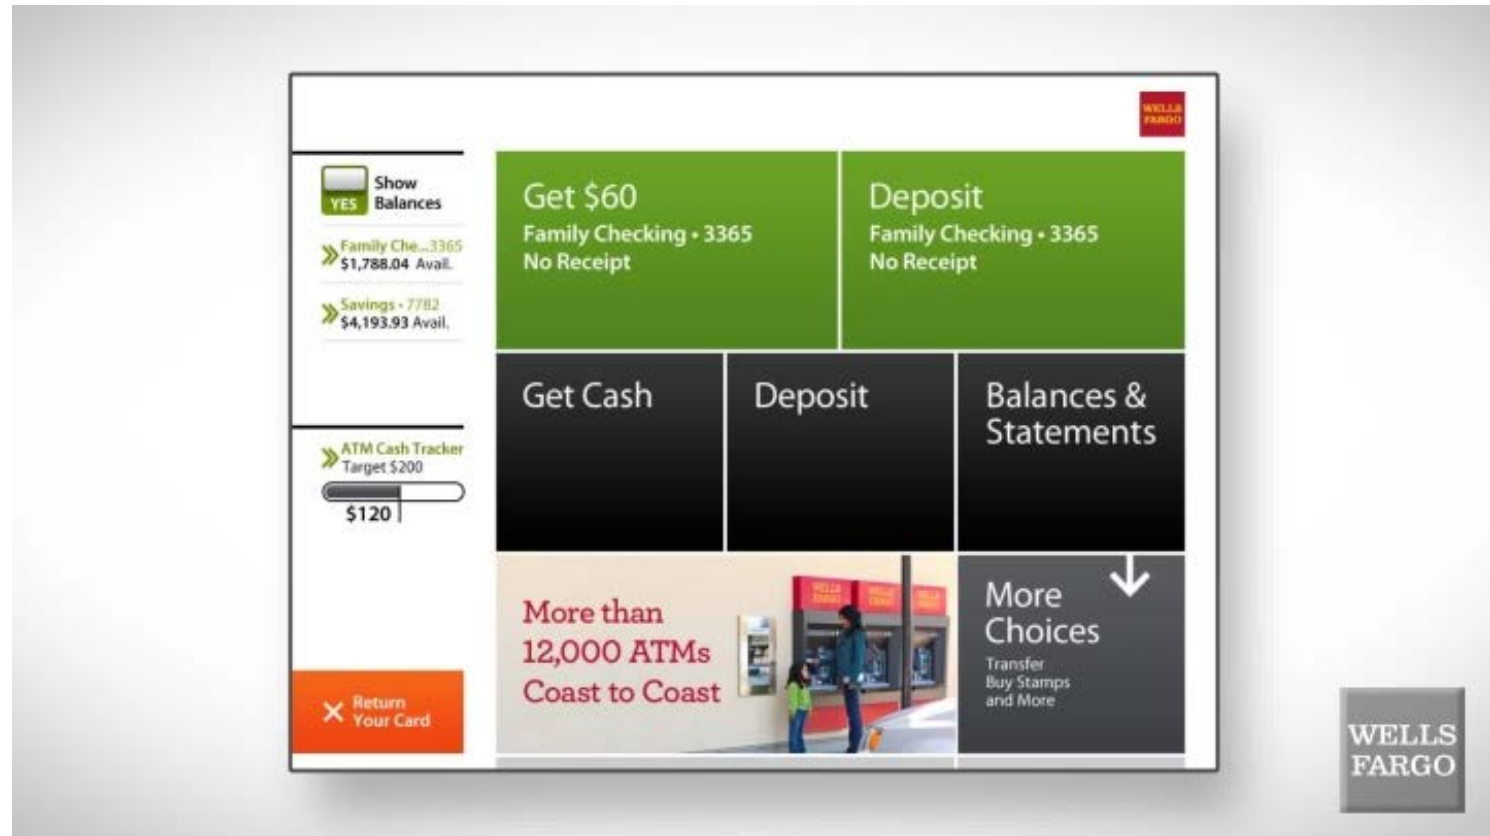
Transactions



Transactions



Motivating example: an ATM



Read Balance
Give money
Update Balance

vs

Read Balance
Update Balance
Give money

It's not just about having the **correct** balance...



Visa does > 60,000 TXNs/sec with users & merchants

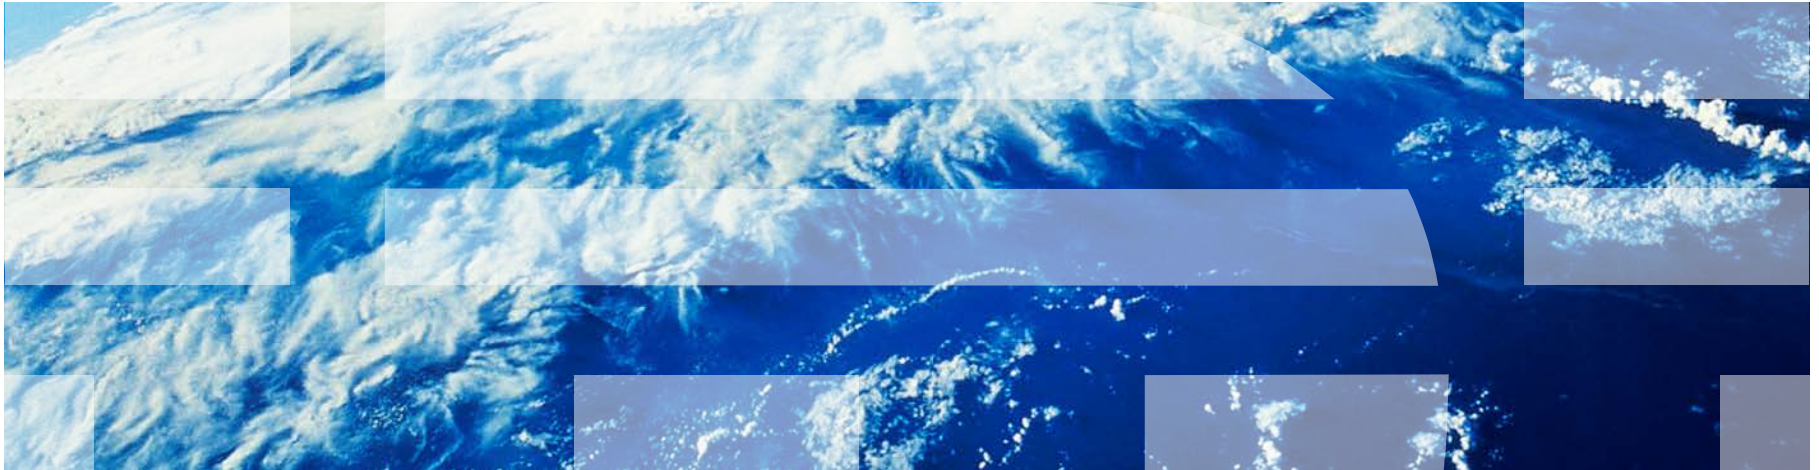
Want your 4\$ Starbucks transaction to wait for a 10k\$ bet in Las Vegas ?
(Transactions can (1) be quick or take a long time, (2) unrelated to you)

Transactions are not just used for finance



- Transactions are at the core of
- payment, stock market, banks, ticketing
 - Gmail, Google Docs (e.g., multiple people editing)

Transactions





Example: monthly bank interest transaction

Money

Account	...	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	...	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22



'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

UPDATE Money
SET Balance = Balance * 1.1



Example: monthly bank interest transaction with crash

Money			
Account	...	Balance (\$)	
3001		500	
4001		100	
5001		20	
6001		60	
3002		80	
4002		-200	
5002		320	
...		...	
30108		-100	
40008		100	
50002		20	

Money (@10:45 am)			
Account	...	Balance (\$)	
3001		550	??
4001		110	
5001		22	
6001		66	
3002		88	
4002		-200	??
5002		320	??
...			
30108		-110	
40008		110	
50002		22	??

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am



Transactions: Basic Definition

A transaction (“TXN”) is a sequence of one or more *operations* (reads or writes) which reflects *a single real-world transition*.

TXN either happened completely or not at all

START TRANSACTION

UPDATE Product

SET Price = Price – 1.99

WHERE pname = ‘Gizmo’

COMMIT



Transactions in SQL

- In “ad-hoc” SQL, each statement = one transaction
- In a program, multiple statements can be grouped together as a transaction

```
START TRANSACTION
```

```
    UPDATE Bank SET amount = amount - 100
```

```
    WHERE name = 'Bob'
```

```
    UPDATE Bank SET amount = amount + 100
```

```
    WHERE name = 'Joe'
```

```
COMMIT
```

Motivation for Transactions

Grouping user actions (reads & writes) into *transactions* helps with two goals:

1. **Recovery and Durability:** Keeping the DB data consistent and durable in the face of crashes, aborts, system shutdowns, etc.
2. **Concurrency:** Achieving better performance by parallelizing TXNs *without* creating anomalies

Motivation -- Recovery & Durability

1. **Recovery and durability** of user data is essential for reliable database (and other data science systems)

- The database may experience crashes (e.g. power outages, etc.)
- Individual TXNs may be aborted (e.g. by the user)

Idea: Make sure that TXNs are either **durably stored in full, or not at all**; keep log to be able to “roll-back” TXNs



Protection against crashes / aborts

Client 1:

```
INSERT INTO SmallProduct(name, price)
SELECT pname, price
FROM Product
WHERE price <= 0.99
```

Crash / abort!

```
DELETE Product
WHERE price <=0.99
```

What goes wrong?

Protection against crashes / aborts

Client 1:

```
START TRANSACTION
INSERT INTO SmallProduct(name, price)
  SELECT pname, price
  FROM Product
  WHERE price <= 0.99

DELETE Product
  WHERE price <=0.99
COMMIT OR ROLLBACK
```

Now we'd be fine! We'll see how / why this lecture

Motivation -- Concurrent execution

2. **Concurrent** execution of user programs is essential for good database performance.

- Disk accesses may be frequent and slow- optimize for throughput (# of TXNs), trade for latency (time for any one TXN)
- Users should still be able to execute TXNs as if in isolation and such that consistency is maintained

Idea: Have the database handle running several user TXNs concurrently, in order to keep throughput high



Multiple users: single statements

Client 1: **UPDATE** Product
 SET Price = Price – 1.99
 WHERE pname = 'Gizmo'

Client 2: **UPDATE** Product
 SET Price = Price*0.5
 WHERE pname = 'Gizmo'

Two managers attempt to discount products *concurrently*-
What could go wrong?

Multiple users: single statements

Client 1: START TRANSACTION

UPDATE Product
SET Price = Price – 1.99
WHERE pname = 'Gizmo'


COMMIT

Client 2: START TRANSACTION

UPDATE Product
SET Price = Price*0.5
WHERE pname='Gizmo'

COMMIT

Lecture 4



Recap of lecture 3

■ **SQL (continued)**

- Foreign keys (adding a constraint on insertions)
- JOINS
 - Inner
 - Outer
- Aggregations: SUM/COUNT/MIN/MAX/AVG
- GROUP BY + aggregates
- HAVING
- Nested queries
 - Breaking into steps
 - Including aggregates, GROUP BY

■ **Transactions motivation:**

- Key idea: either entire operations is executed, or it is not
- Transactions in SQL
- Two key challenges
 - Durability and recovery (surviving crashes, making sure we can rollback)
 - Concurrent operations (allow transactions to execute in parallel and in isolation)

Today: deep dive on transactions

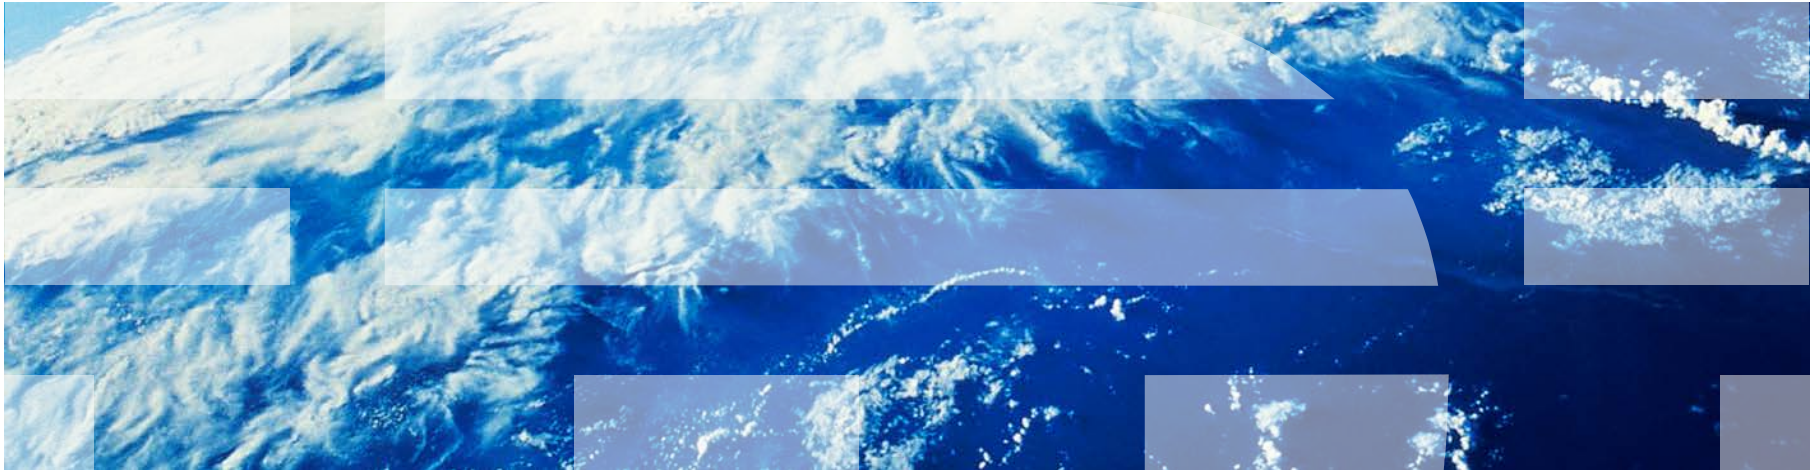
- ACID
- Atomicity and durability via logging
 - Write ahead log
- Consistency and isolation via locking
 - 2 Phase Locking
- If we have time: conflict serializability

Logistics

- Written homework will be released at the end of this week
 - Meant as preparation for midterm
 - After this class, you will have the material to solve most (but not all) of it
- Reminder: SQL programming homework due next week at **noon on Thursday**
 - We covered all the material to solve it

ACID

Atomicity, Consistency, Isolation, Durability



Transaction Properties: ACID

- **A**tomic
 - State shows either all the effects of txn, or none of them
- **C**onsistent
 - Txn moves from a state where integrity holds, to another where integrity holds
- **I**solated
 - Effect of txns is the same as txns running one after another (ie looks like batch mode)
- **D**urable
 - Once a txn has committed, its effects remain in the database

ACID continues to be a source of great debate!

ACID: Atomicity

- TXN's activities are atomic: all or nothing
 - Intuitively: in the real world, a transaction is something that would either occur *completely* or *not at all*
- Two possible outcomes for a TXN
 - It *commits*: all the changes are made
 - It *aborts*: no changes are made

ACID: Consistency

- The tables must always satisfy user-specified *integrity constraints*
 - *Examples*:
 - Account number is unique
 - Stock amount can't be negative
 - Sum of *debits* and of *credits* is 0
- How consistency is achieved:
 - Programmer **writes** a **TXN** to go from one consistent state to a consistent state
 - *System* makes sure that the **TXN** is atomic
 - → Assuming system maintaining atomicity, this is often the user's responsibility

ACID: Isolation

- A transaction executes concurrently with other transactions
- **Isolation**: the effect is as if each transaction executes in *isolation* of the others.
- E.g. Should not be able to observe changes from other transactions during the run

ACID: Durability

- The effect of a TXN must continue to exist (*“persist”*) after the TXN
 - And after the whole program has terminated
 - And even if there are power failures, crashes, etc.
 - And etc...
- Means: Write data to disk
 - And in data center settings: replicate data, backup, etc.



Challenges for ACID properties

- In spite of power failures (i.e., in spite of loss of memory)
- Users may abort the program: need to “rollback changes”
 - Need to *log* what happened
- Many users executing concurrently


And all this with... Scalability and/or Performance!!

A Note: ACID is contentious!

- Many debates over ACID, both **historically** and **currently**
- Some “NoSQL” DBs relax ACID
- In turn, now “NewSQL” reintroduces ACID compliance to NoSQL-style DBs...



Atomicity and Durability via Logging



Goal for this lecture: Ensuring Atomicity & Durability

ACID

- Atomicity:
 - TXNs should either happen completely or not at all
 - If abort / crash during TXN, *no* effects should be seen

TXN 1



Crash / abort

No changes
persisted

- Durability:
 - If DB stops running, changes due to completed TXNs should all persist
 - *Just store on stable disk*

TXN 2



All changes
persisted

We'll focus on how to accomplish atomicity (via logging)

Basic Idea: (Physical) Logging

Idea:

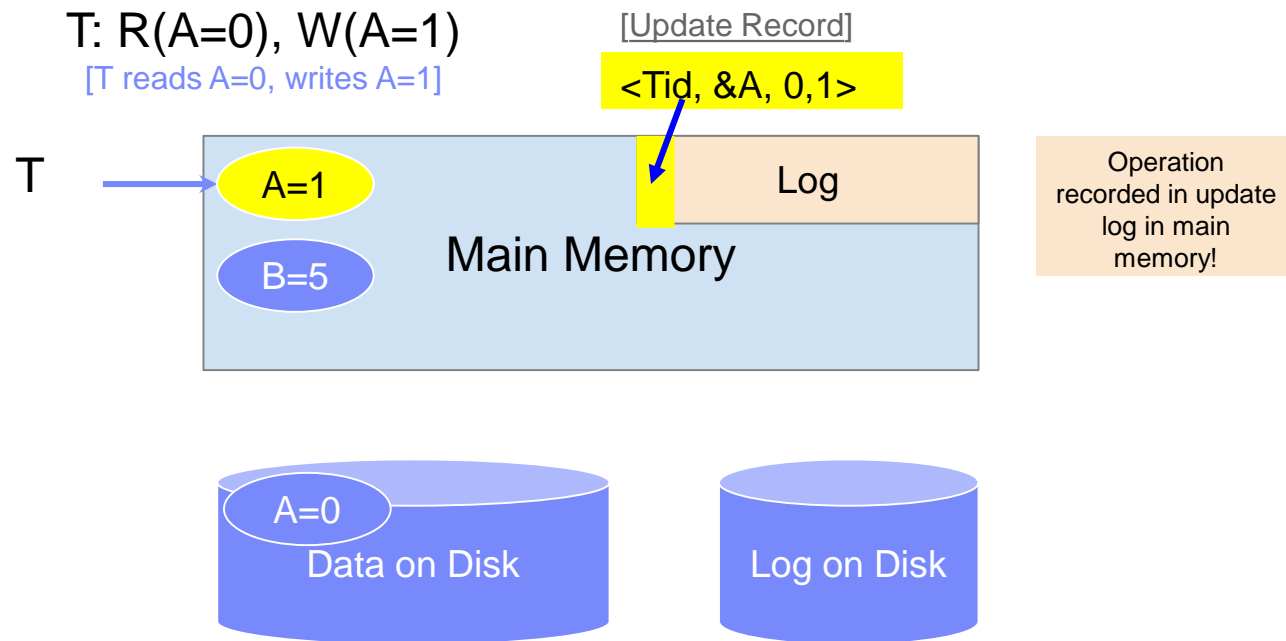
- Log consists of an ordered list of update records
- Log record contains UNDO information for every update!
<TransactionID, location, old data, new data>

What DB does?

- Owns the log “service” for all applications/transactions.
- Transparent to application or transaction
- Sequential writes to log, can **flush** — force writes to disk

This is sufficient to UNDO any transaction!

A picture of logging



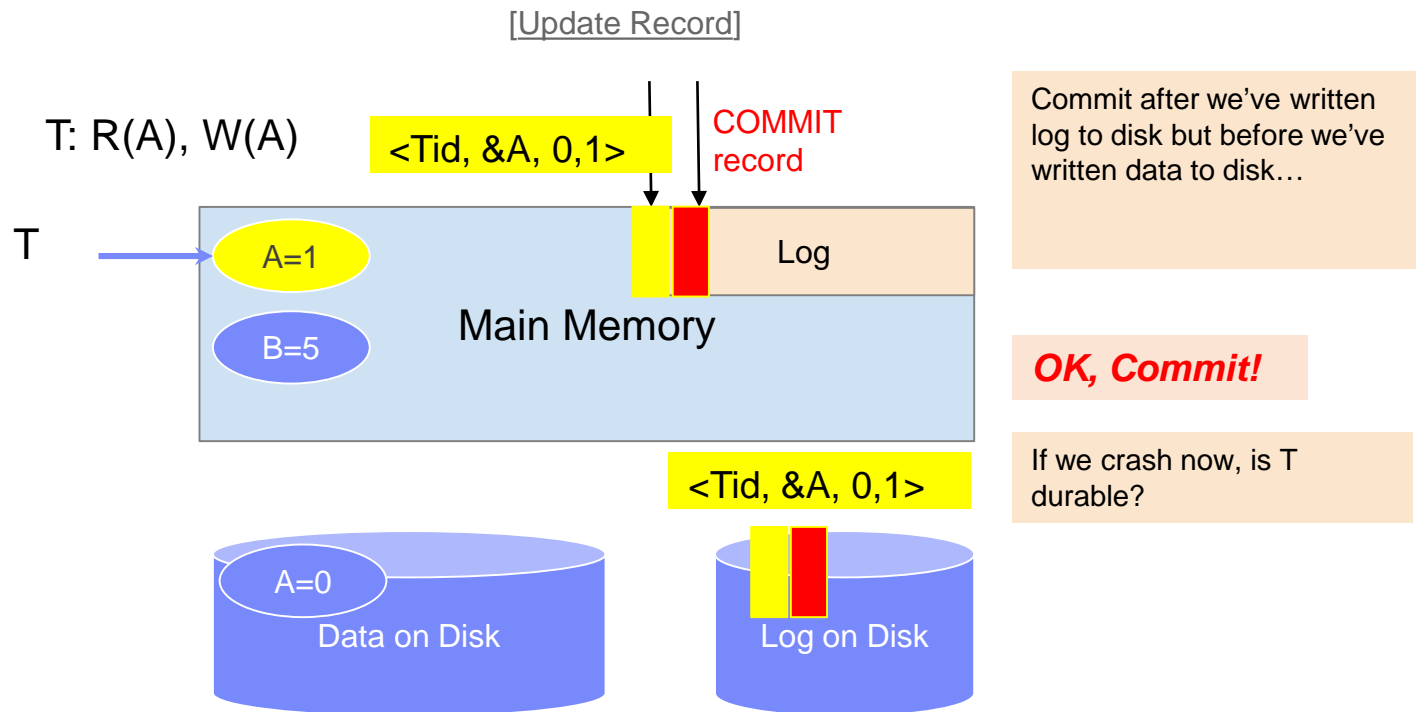


Why do we need logging for atomicity?

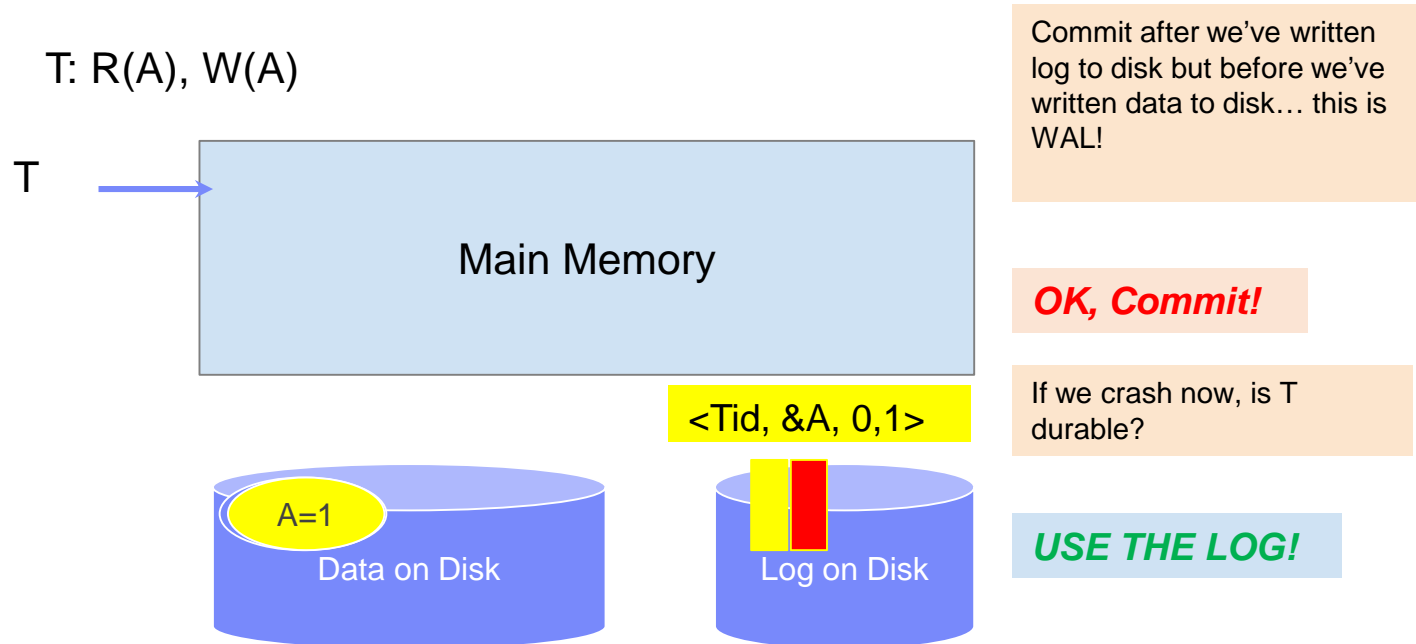
- Couldn't we just write TXN to disk **only** once whole TXN complete?
 - Then, if abort / crash and TXN not complete, it has no effect- atomicity!
 - *With unlimited memory and time, this could work...*
- However, we **need to log partial results of TXNs** because of:
 - Memory constraints (enough space for full TXN??)
 - Time constraints (what if one TXN takes very long?)

We need to write partial results to disk!
...And so we need a **log** to be able to **undo** these partial results!

Write-ahead Logging (WAL) Commit Protocol



Write-ahead Logging (WAL) Commit Protocol



Write-Ahead Logging (WAL)

Algorithm: WAL

For each record update, write Update Record into LOG

Follow two **Flush** rules for LOG

- **Rule1:** **Flush** Update Record into LOG before corresponding data page goes to storage
- **Rule2:** Before TXN commits,
 - **Flush** all Update Records to LOG
 - **Flush** COMMIT Record to LOG

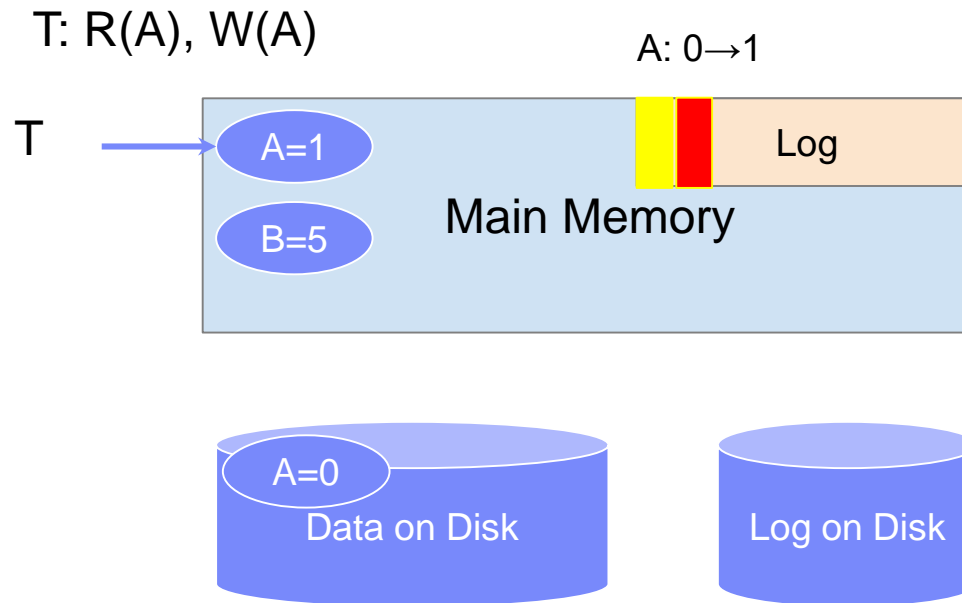
→ **Durability**

→ **Atomicity**

Transaction is committed *once COMMIT record is on stable storage*



Incorrect Commit Protocol #1



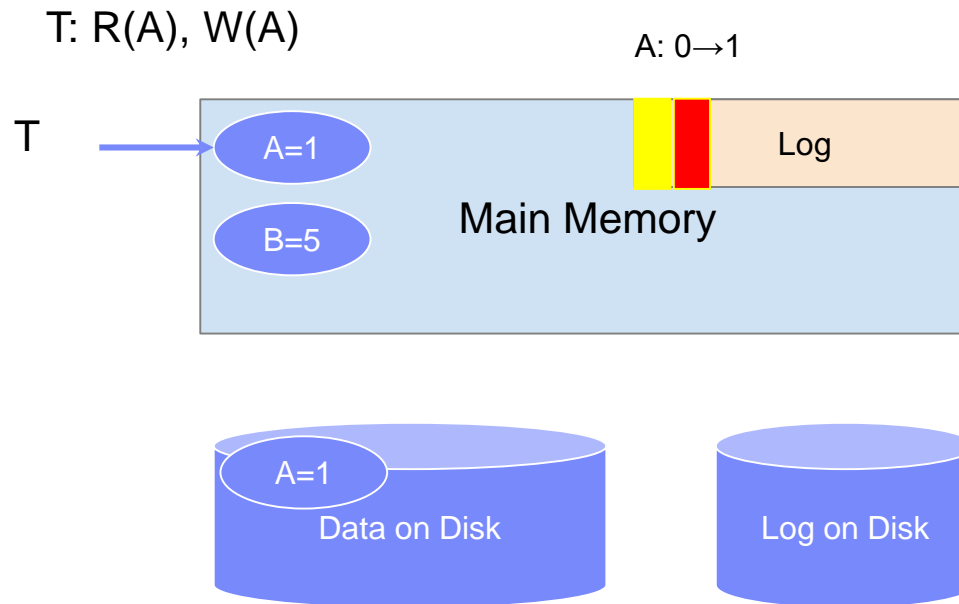
Let's try committing
before we've written
either data or log to
disk...

OK, Commit!

If we crash now, is T
durable?

Lost T's update!

Incorrect Commit Protocol #2



Let's try committing *after* we've written data but *before* we've written log to disk...

OK, Commit!

If we crash now, is T durable? Yes! Except...

How do we know whether T was committed??

Bank interest example: full run

Money			Money (@4:29 am day+1)			WAL (@4:29 am day+1)			
Account	...	Balance (\$)	Account	...	Balance (\$)				
3001		500	3001		550	T-Monthly-423	START TRANSACTION		
4001		100	4001		110	T-Monthly-423	3001	500	550
5001		20	5001		22	T-Monthly-423	4001	100	110
6001		60	6001		66	T-Monthly-423	5001	20	22
3002		80	3002		88	T-Monthly-423	6001	60	66
4002		-200	4002		-220	T-Monthly-423	3002	80	88
5002		320	5002		352	T-Monthly-423	4002	-200	-220
...		T-Monthly-423	5002	320	352
30108		-100	30108		-110	T-Monthly-423
40008		100	40008		110	T-Monthly-423	30108	-100	-110
50002		20	50002		22	T-Monthly-423	40008	100	110
						T-Monthly-423	50002	20	22
						T-Monthly-423	COMMIT		

Update
Records

Commit
Record

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

START TRANSACTION

UPDATE Money

SET Amt = Amt * 1.10

COMMIT

Bank interest example: with crash

Money			Money (@10:45 am)			WAL log (@10:29 am)			
Account	...	Balance (\$)	Account	...	Balance (\$)				
3001		500	3001		550	??	T-Monthly-423	START TRANSACTION	
4001		100	4001		110		T-Monthly-423	3001	500 550
5001		20	5001		22		T-Monthly-423	4001	100 110
6001		60	6001		66		T-Monthly-423	5001	20 22
3002		80	3002		88		T-Monthly-423	6001	60 66
4002		-200	4002		-200	??	T-Monthly-423	3002	80 88
5002		320	5002		320	??	T-Monthly-423
...			T-Monthly-423	30108	-100 -110
30108		-100	30108		-110		T-Monthly-423	40008	100 110
40008		100	40008		110		T-Monthly-423	50002	20 22
50002		20	50002		22	??	T-Monthly-423	4002	-200 -220
							T-Monthly-423	5002	320 352

'T-Monthly-423'

Monthly Interest 10%
4:28 am Starts run on 10M bank accounts
Takes 24 hours to run

Network outage at 10:29 am,
System access at 10:45 am

Did T-Monthly-423 complete?
Which tuples are bad?

Case1: T-Monthly-423 was crashed
Case2: T-Monthly-423 completed. 4002
deposited 20\$ at 10:45 am

Bank interest example: with recovery

Money (@10:45 am)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-200
5002		320
...		
30108		-110
40008		110
50002		22

Money (after recovery)

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

WAL log (@10:29 am)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22

System recovery (after 10:45 am)

1. Rollback uncommitted transactions
 - Restore old values from WALlog (if any)
 - Notify developers about aborted txn
1. Redo Recent transactions (w/ new values)
2. Back in business; Redo (any pending) transactions

A word on performance

- Question: why is a WAL a good idea?
- Answer: updates to WAL are in sequential order!
- Recall: sequential writes are very important both for flash and magnetic disk
 - In a couple of lectures we will understand why

An example of why sequential writes matter

Money

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22

WAL (@4:29 am day+1)

T-Monthly-423	START TRANSACTION		
T-Monthly-423	3001	500	550
T-Monthly-423	4001	100	110
T-Monthly-423	5001	20	22
T-Monthly-423	6001	60	66
T-Monthly-423	3002	80	88
T-Monthly-423	4002	-200	-220
T-Monthly-423	5002	320	352
T-Monthly-423
T-Monthly-423	30108	-100	-110
T-Monthly-423	40008	100	110
T-Monthly-423	50002	20	22
T-Monthly-423	COMMIT		

Cost to update all data

10M bank accounts → 10M individual random writes? (worst case)


(@10 ms per write for magnetic disk, that's 100,000 secs)



Speedup for commit

100,000 secs vs 1 sec when written sequentially!!!

Concurrency and Locking for Transactions



Back to our bank example

Money

Account	Balance (\$)
3001		500
4001		100
5001		20
6001		60
3002		80
4002		-200
5002		320
...		...
30108		-100
40008		100
50002		20

Money (@4:29 am day+1)

Account	Balance (\$)
3001		550
4001		110
5001		22
6001		66
3002		88
4002		-220
5002		352
...		...
30108		-110
40008		110
50002		22



Other Transactions

10:02 am Acct 3001: Wants 600\$

11:45 am Acct 5002: Wire for 1000\$

.....

.....

2:02 pm Acct 3001: Debit card for \$12.37

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

UPDATE Money
SET Balance = Balance * 1.1

Q: How do I not wait for a day to access my \$\$\$s?

Big idea: locks

■ Intuition

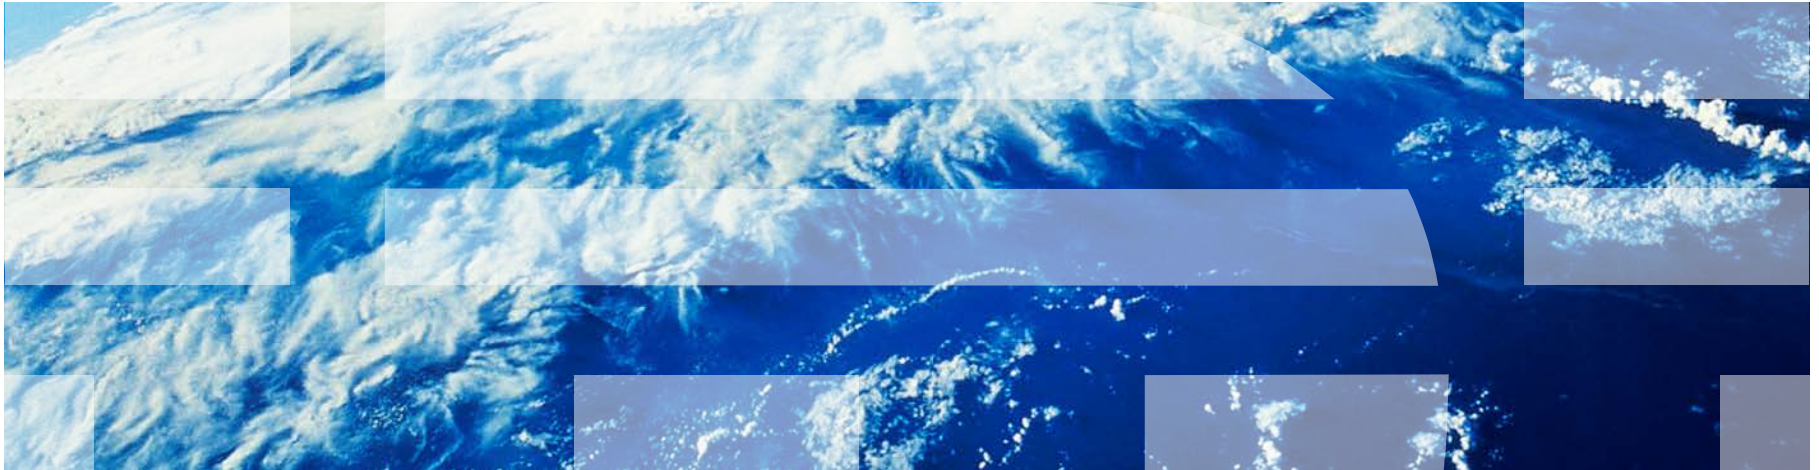
- "Lock" each record for shortest time possible

■ Key questions

- Which records?
- For how long?
- What is the algorithm for holding them?



Concurrency, Scheduling and Anomalies



Concurrency: Isolation & Consistency

- DB is responsible for concurrency so that...

Isolation is maintained: Users must be able to execute each TXN as if they were the only user

ACID

Consistency is maintained: TXNs must leave the DB in consistent state

ACCID



Example- consider two TXNs:

T1: START TRANSACTION

UPDATE Accounts

SET Amt = Amt + 100

WHERE Name = 'A'

UPDATE Accounts

SET Amt = Amt - 100

WHERE Name = 'B'

COMMIT

T1 transfers \$100 from B's account to A's account

T2: START TRANSACTION

UPDATE Accounts

SET Amt = Amt * 1.06

COMMIT

T2 credits both accounts with a 6% interest payment

Note:

1. DB does not care if T1 \rightarrow T2 or T2 \rightarrow T1 (which TXN executes first)
2. If developer does, what can they do? (Put T1 and T2 inside 1 TXN)



Example

T_1

$A += 100$

$B -= 100$

T1 transfers \$100 from B's account to A's account

T_2

$A *= 1.06$

$B *= 1.06$

T2 credits both accounts with a 6% interest payment

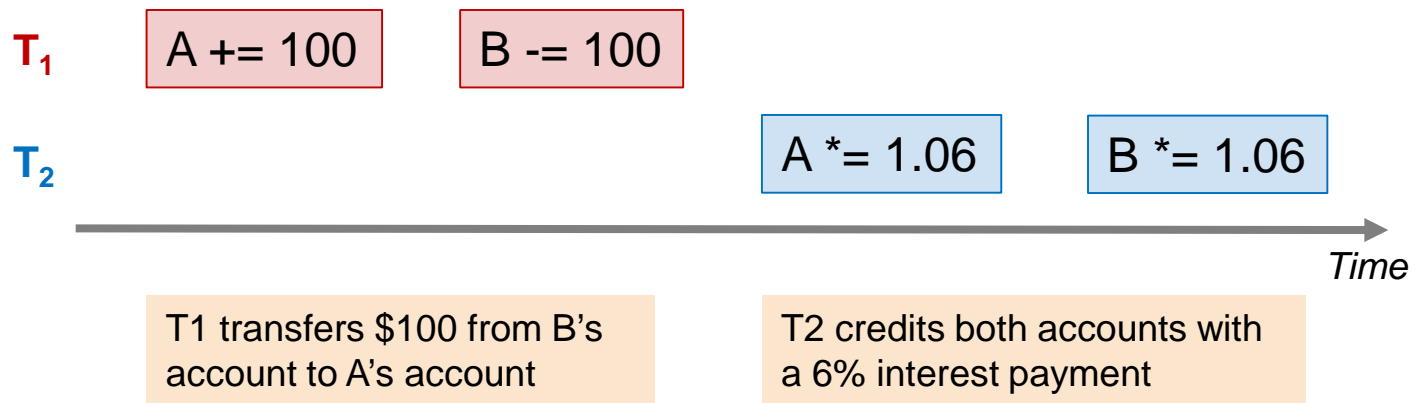
Goal for scheduling transactions:

- Interleave transactions to boost performance
- Data stays in a good state after commits and/or aborts (ACID)



Example- consider two TXNs:

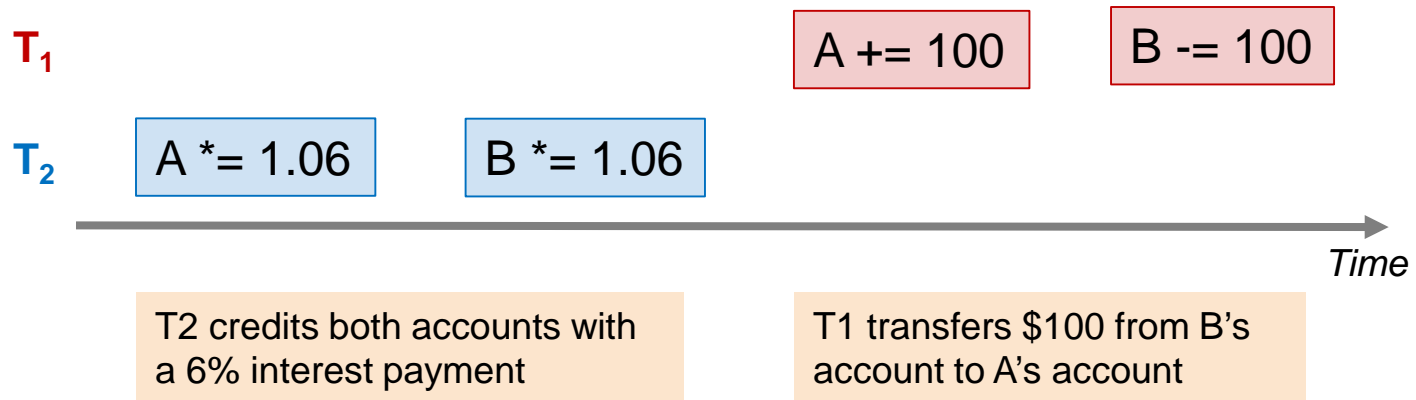
We can look at the TXNs in a timeline view- serial execution:





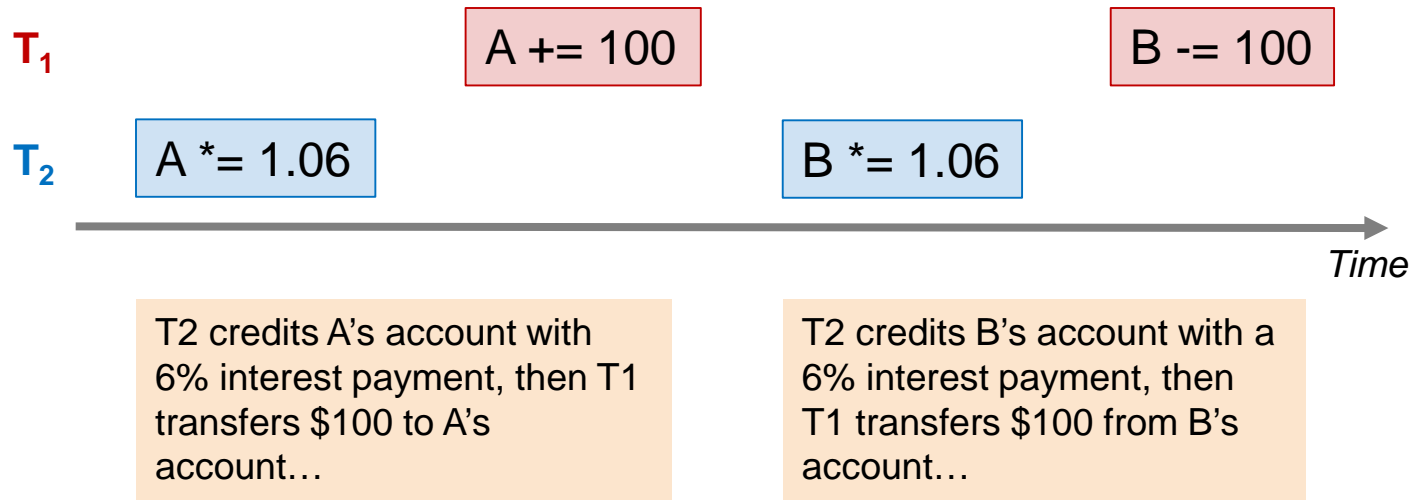
Example- consider two TXNs:

The TXNs could occur in either order... DB allows!



Example- consider two TXNs:

The DB can also **interleave** the TXNs



Interleaving & Isolation

- The DB has freedom to interleave TXNs
- However, it must pick an interleaving or schedule such that isolation and consistency are maintained
- ⇒ Must be *as if* the TXNs had executed serially!

“With great power
comes great
responsibility”

ACID

DB must pick a schedule which maintains isolation
& consistency

Same
result!

Scheduling examples

*Starting
Balance*

A	B
\$50	\$200

Serial schedule T_1, T_2 :

T₁	A += 100	B -= 100
----------------------	----------	----------

T_2	$A^* = 1.06$	$B^* = 1.06$
-------	--------------	--------------

A	B
\$159	\$106

Interleaved schedule B:

T₁ A += 100 B -= 100

T_2	$A^* = 1.06$	$B^* = 1.06$
-------	--------------	--------------

A	B
\$159	\$112

Different
result than
serial
 T_1, T_2 !

Scheduling examples

Starting
Balance

A	B
\$50	\$200

Serial schedule T_2, T_1 :

T_1

A += 100 B -= 100

T_2

A *= 1.06 B *= 1.06

A	B
\$153	\$112

Interleaved schedule B:

T_1

A += 100

B -= 100

T_2


A *= 1.06 B *= 1.06

A	B
\$159	\$112

Different
result than
serial
 T_2, T_1
ALSO!

Scheduling examples

Interleaved schedule B:



This schedule is different than ***any serial order!*** We say that it is **not serializable**



Scheduling Definitions

- A **serial schedule** is one that does not interleave the actions of different transactions
- A and B are **equivalent schedules** if, *for any database state*, the effect on DB of executing A **is identical to** the effect of executing B
- A **serializable schedule** is a schedule that is equivalent to ***some*** serial execution of the transactions.





The word “**some**” makes this definition powerful & tricky!



Serial Schedules




Interleaved Schedules



Serial Schedules	S1, S2
Serializable Schedules	S3, S4 (And S1, S2)
Equivalent Schedules	<S1, S3> <S2, S4>
Non-serializable (Bad) Schedules	S5, S6


Conflicts and Anomalies




General DB model:

Concurrency as Interleaving TXNs

Serial Schedule



Interleaved Schedule



Each action in the TXNs
*reads a value and then
writes one back*

For our purposes, having TXNs
occur concurrently means
**interleaving their component
actions (R/W)**

We call the particular order
of interleaving a **schedule**



Conflict Types

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write

Thus, there are three types of conflicts:

- Read-Write conflicts (RW)
- Write-Read conflicts (WR)
- Write-Write conflicts (WW)


Why no “RR Conflict”?

Note: **conflicts** happen often in many real world transactions. (E.g., two people trying to book an airline ticket)



Conflicts

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write



All "conflicts"!

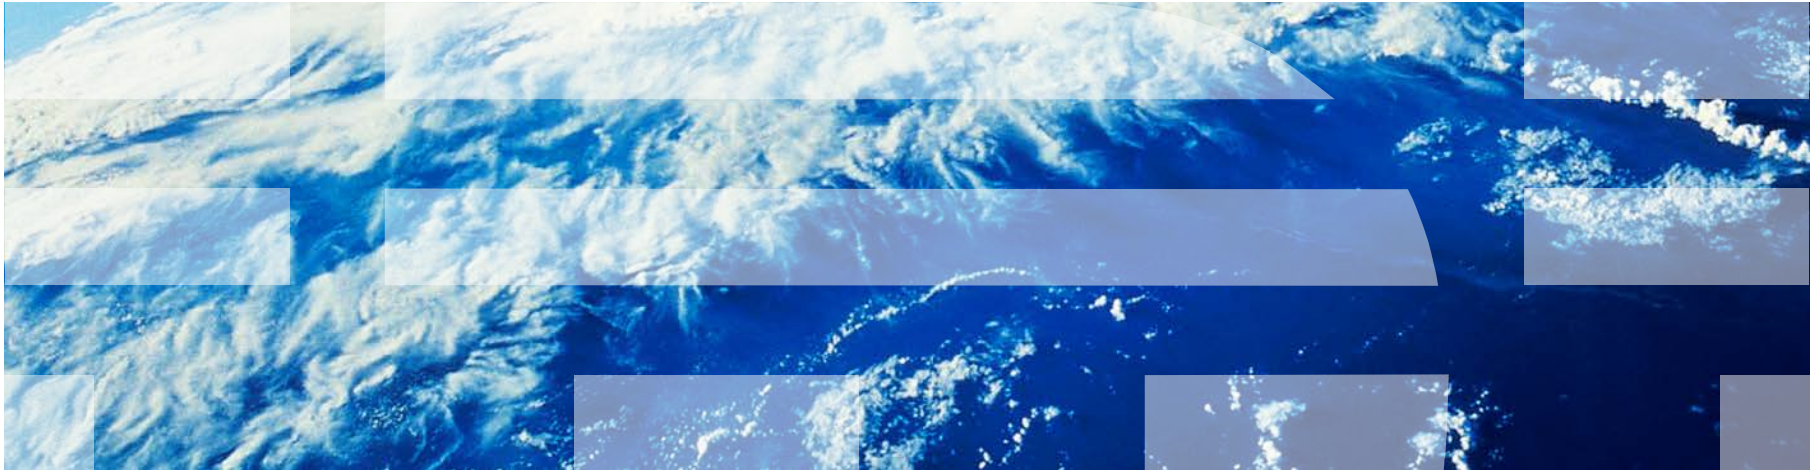
Note: Conflicts vs. Anomalies

Conflicts are in both “good” and “bad” schedules
(they are a property of transactions)

Goal: Avoid Anomalies while interleaving transactions with conflicts!

- Do not create “bad” schedules where isolation and/or consistency is broken (i.e., Anomalies)

Conflict Serializability



Conflict Serializability

Two schedules are **conflict equivalent** if:

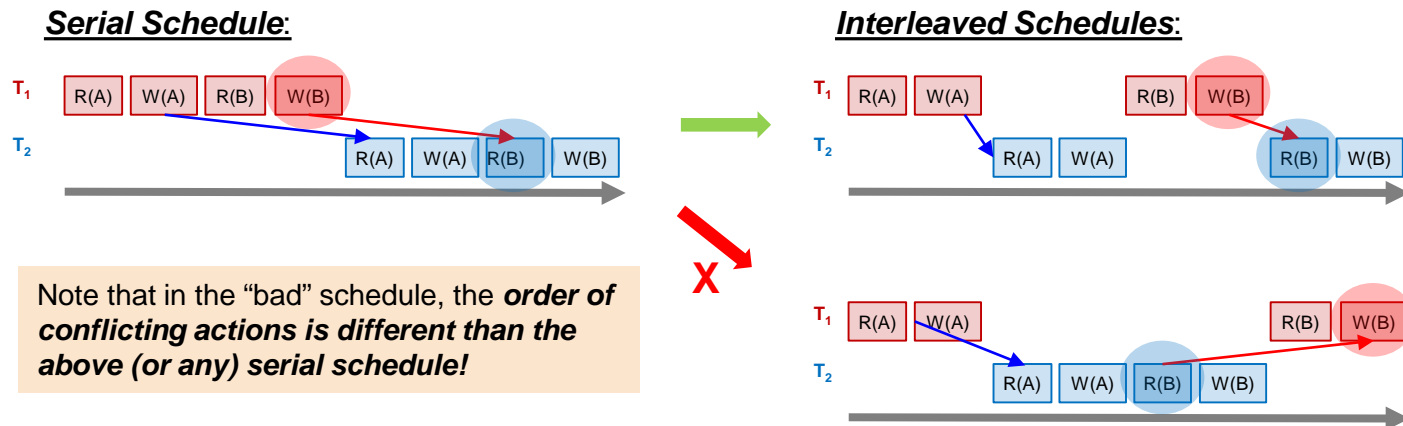
- They involve *the same actions of the same TXNs*
- Every *pair of conflicting actions* of two TXNs are *ordered in the same way*

Schedule S is **conflict serializable** if S is *conflict equivalent* to some serial schedule

Conflict serializable \Rightarrow serializable

So if we have conflict serializable, we have consistency & isolation!


Example “Good” vs. “bad” schedules



Conflict serializability provides us with an operative notion of “good” vs. “bad” schedules! “Bad” schedules create data Anomalies


The Conflict Graph

- Let's now consider looking at conflicts **at the TXN level**
- Consider a graph where the **nodes are TXNs**, and there is an edge from $T_i \rightarrow T_j$ **if any actions in T_i precede and conflict with any actions in T_j**





What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:



Interleaved Schedules:




What can we say about “good” vs. “bad” conflict graphs?

Serial Schedule:



Simple!

Interleaved Schedules:

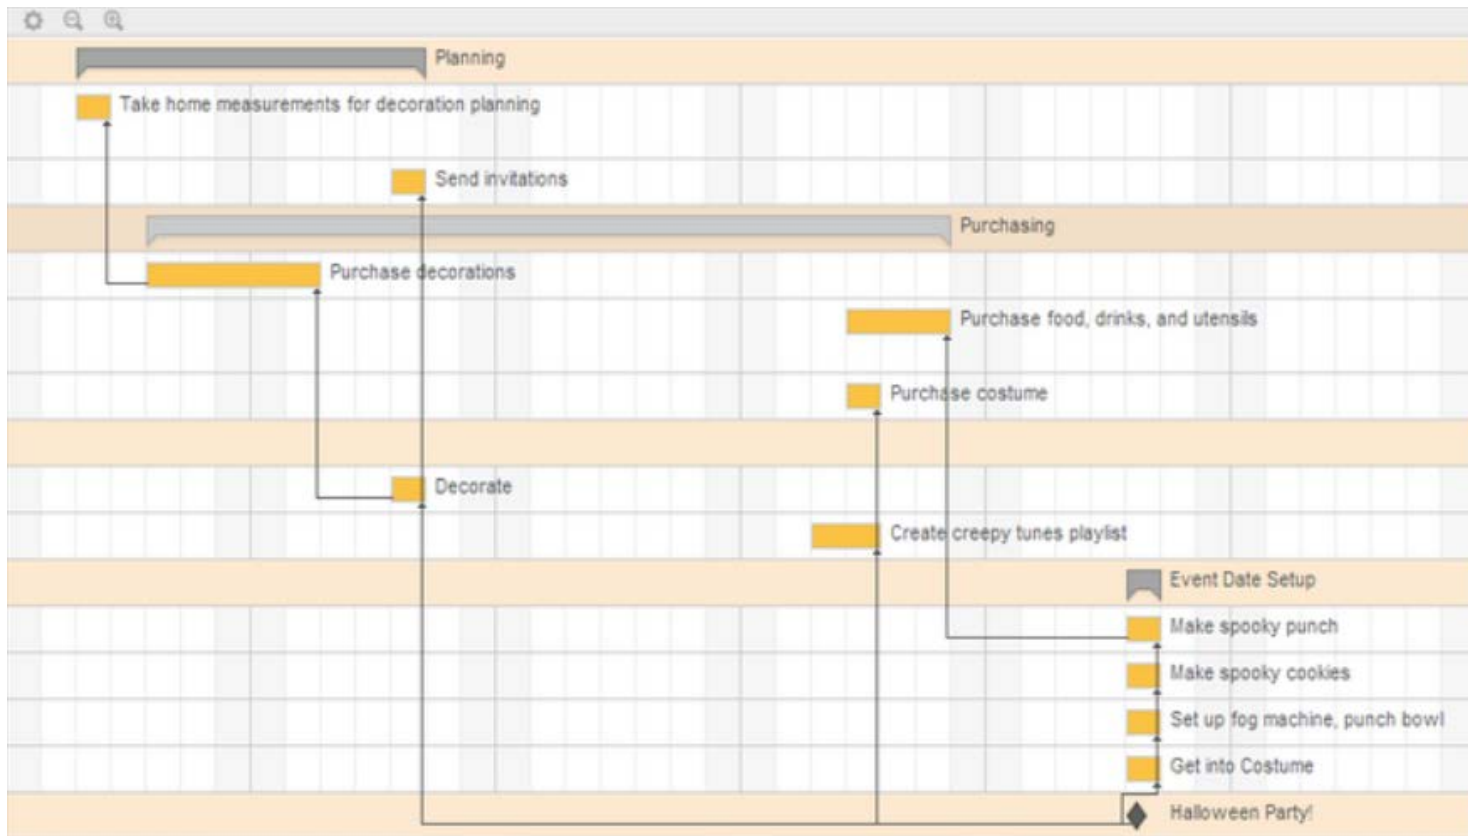


Theorem: Schedule is **conflict serializable** if and only if its conflict graph is **acyclic**

DAGs & Topological Orderings

- A **topological ordering** of a directed graph is a linear ordering of its vertices that respects all the directed edges
- A directed **acyclic** graph (DAG) always has one or more **topological orderings**
 - (And there exists a topological ordering *if and only if* there are no directed cycles)

Example: Project dependencies




How would you plan?
What if there are cycles? (dependencies)



DAGs & Topological Orderings


- Ex: What is one possible topological ordering here?



Ex: 0, 1, 2, 3 (or: 0, 1, 3, 2)

DAGs & Topological Orderings

- Ex: What is one possible topological ordering here?



There is none!

Connection to conflict serializability

- In the conflict graph, a topological ordering of nodes corresponds to **a serial ordering of TXNs**
- Thus an **acyclic** conflict graph → conflict serializable!

Theorem: Schedule is **conflict serializable** if and only if its conflict graph is **acyclic**

Example with 5 transactions

Schedule S1


Good or Bad schedule?
Conflict serializable?

w1(A)	r2(A)	w1(B)	w3(C)	r2(C)	r4(B)	w2(D)	w4(E)	r5(D)	w5(E)
-------	-------	-------	-------	-------	-------	-------	-------	-------	-------

Step1
Find conflicts
(RW, WW, WR)

T1	w1(A)		w1(B)						
T2		r2(A)			r2(C)		w2(D)		
T3				w3(C)					
T4						r4(B)		w4(E)	
T5								r5(D)	w5(E)

Step2
Build Conflict graph
Acyclic?



Acyclic
⇒ Conflict serializable!
⇒ Serializable

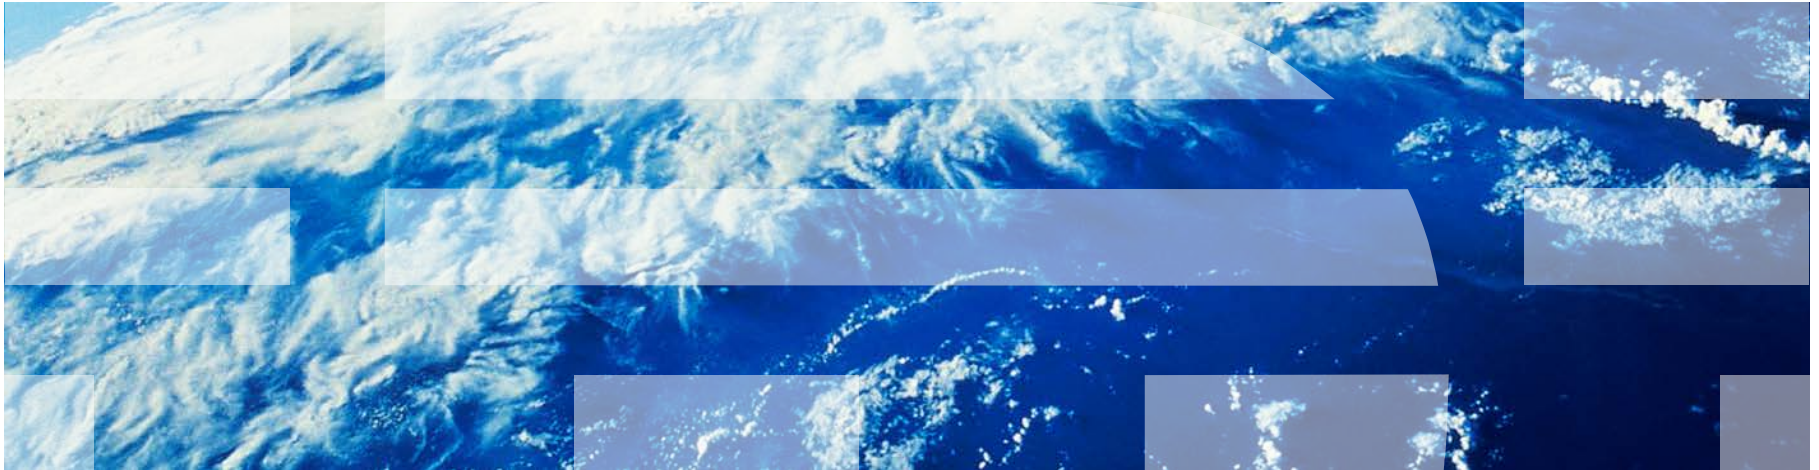
Step3
Example serial schedule
Conflict Equiv to S1

S2									
T3	T1	T1	T4	T4	T2	T2	T2	T5	T5
w3(C)	w1(A)	w1(B)	r4(B)	w4(E)	r2(A)	r2(C)	w2(D)	r5(D)	w5(E)

Summary

- Concurrency achieved by **interleaving TXNs** such that **isolation & consistency** are maintained
 - We formalized a notion of **serializability** that captured such a “good” interleaving schedule
- We defined **conflict serializability**

2PL: A Simple Locking Algorithm



Strict Two-Phase Locking (2PL)

- Algorithm: *strict two-phase locking* - as a way to deal with concurrency
 - Guarantees conflict serializability
 - (if it completes- see upcoming...)
- Also (*conceptually*) straightforward to implement, and transparent to the user!

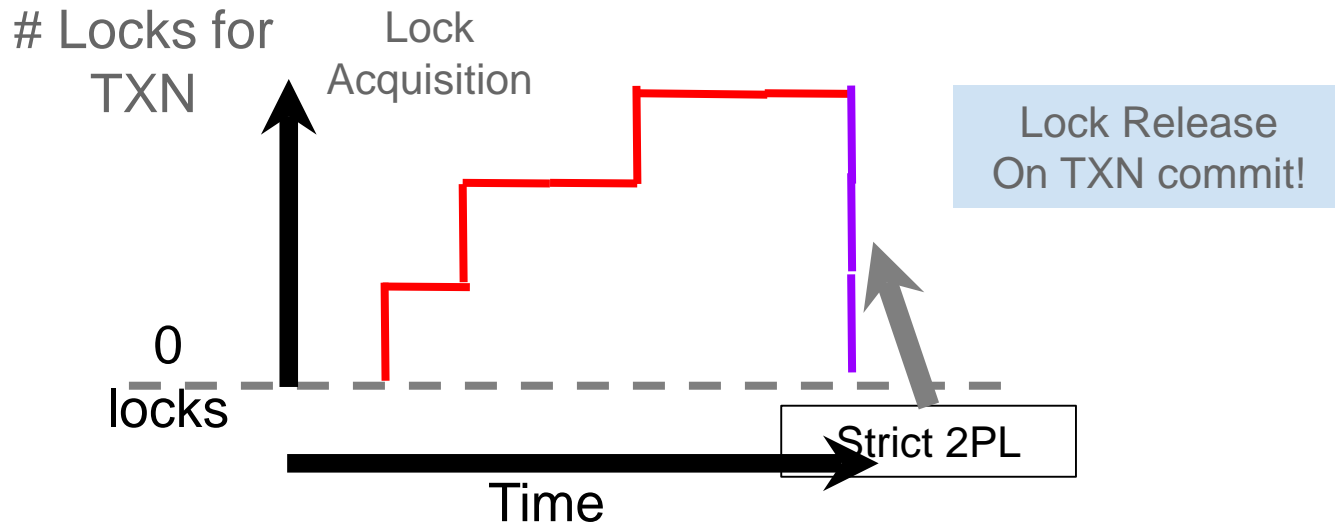
Strict Two-phase Locking (2PL) Protocol

TXNs obtain:

- An **X (*exclusive*) lock** on object before **writing**.
 - If a TXN holds, no other TXN can get a lock (S or X) on that object.
- An **S (*shared*) lock** on object before **reading**
 - If a TXN holds, no other TXN can get an X lock on that object
- All locks held by a TXN are released when TXN completes.

Note: Terminology here- “exclusive”, “shared”- meant to be intuitive- no tricks!

Picture of 2-Phase Locking (2PL)



2PL: A transaction can not request additional locks once it releases any locks. Thus, there is a “growing phase” followed by a “shrinking phase”.

Strict 2PL: Release locks only at COMMIT (COMMIT Record flushed) or ABORT

Strict 2PL

If a schedule follows strict 2PL, it is **conflict serializable**...

- ...and thus serializable
- ...and we get isolation & consistency!

Popular implementation

- Simple !
- Produces subset of *all* conflict serializable schedules
- There are MANY more complex LOCKING schemes with better performance. (See CS Database classes)
- One key, subtle problem (next)

Deadlock Detection

T_1


$S(A)$	$R(A)$
--------	--------

T_2



First, T_1 requests a shared lock on A to read from it

Deadlock Detection



Next, T_2 requests a shared lock on B to read from it


Deadlock Detection: Example




T₂ then requests an exclusive lock on A to write to it- **now T₂ is waiting on T₁...**

Waits-For graph: Track which Transactions are waiting
IMPORTANT: WAITS-FOR graph different than
CONFLICT graph we learnt earlier !

Deadlock Detection: Example



Waits-for graph:



Cycle =
DEADLOCK

Finally, T_1 requests an exclusive lock on B to write to it- **now T_1 is waiting on T_2 ... DEADLOCK!**

Deadlocks

Deadlock: Cycle of transactions waiting for locks to be released by each other.

Two ways of dealing with deadlocks:

- Deadlock prevention

- Deadlock detection

Deadlock Prevention

Conservative 2 Phase Locking (C2PL)

- Obtains all locks before the transaction begins
- If cannot obtain ALL locks, release and try again
- Ensures that no deadlocks occurs
- BUT: can degrade performance

Deadlock Detection

Create the **waits-for graph**:



- Nodes are transactions
- There is an edge from $T_i \rightarrow T_j$ if T_i is *waiting for T_j to release a lock*

Periodically check for (***and break***) cycles in the waits-for graph

Example with 5 Transactions (2PL)

Schedule S1

Execute with 2PL




Waits- For Graph

Example with 5 Transactions (2PL)

Schedule S1

Execute with 2PL



Step 0

X (A)

Step 1

w1(A)

Req S(A)

Step 2

X (B)

Unl B, A

Step 3

Get S(A)

r2(A)

Step 4

X (C)

w3(C)

Unl C

Step 5

S(C)

r2(C)

Step 6

S(B)

r4(B)

Step 7

X(D)

w2(D)

Unl A, C, D

Step 8

X(E)

w4(E)

Unl B, E

Step 9

S(D)

r5(D)

Step 10

X (E)

w5(E)

Unl D, E



Waits- For Graph

Summary

Locking allows only conflict serializable schedules

- If the schedule completes... (it may deadlock!)